

# RIASEC ANALYSIS

```
In [1]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
```

```
In [2]: #df = pd.read_csv("RIASEC.csv")
```

## Reading and Cleaning

```
In [3]: df = pd.read_csv("RIASEC.csv", sep=None, engine="python")

print(df.shape)
print(df.head())
```

```
(8855, 55)
  implementation  R1  R2  R3  R4  R5  R6  R7  R8  I1  ...  C5  C6  C7  C8  \
0              2   3   1   4   2   1   2   1   1   5  ...   2   1   1   2
1              2   1   1   1   1   1   1   1   1   4  ...   1   1   1   1
2              2   3   2   1   1   1   1   2   1   5  ...   3   4   4   4
3              2   3   2   1   2   2   3   1   2   5  ...   1   3   2   1
4              2  -1   2   3   2   3   2   1   3   5  ...   4   3   3   3

  accuracy  elapse  country  fromsearch  age  gender
0        90    222      PT          0    -1    -1
1       100    102      US          0    -1    -1
2        95    264      US          1    -1    -1
3        60    189      SG          0    -1    -1
4        90    197      US          0    -1    -1
```

[5 rows x 55 columns]

```
In [4]: cols = [f"R{i}" for i in range(1, 9)]
print(cols)
R_data = df[cols]
R_data.shape
```

['R1', 'R2', 'R3', 'R4', 'R5', 'R6', 'R7', 'R8']

Out[4]: (8855, 8)

```
In [5]: df = R_data[(R_data != -1).all(axis=1)] # cleans the data and gets rid of rows w
print("Shape after cleaning:", df.shape)
```

Shape after cleaning: (8478, 8)

```
In [6]: # the entries have reduced from 8855 people to 8478 meaning 377 people had inval
```

```
In [7]: 8855-8478
```

Out[7]: 377

```
In [8]: df
```

Out[8]:

	R1	R2	R3	R4	R5	R6	R7	R8
0	3	1	4	2	1	2	1	1
1	1	1	1	1	1	1	1	1
2	3	2	1	1	1	1	2	1
3	3	2	1	2	2	3	1	2
5	3	1	3	4	3	4	3	3
...	...	...	...	...	...	...	...	...
8849	3	3	1	4	2	2	2	1
8851	3	2	3	4	3	2	2	2
8852	4	3	3	3	2	2	1	2
8853	4	4	3	5	4	5	3	4
8854	4	2	4	4	1	4	2	3

8478 rows × 8 columns

In [9]: `df["R1"]`

Out[9]:

```

0      3
1      1
2      3
3      3
5      3
..
8849   3
8851   3
8852   4
8853   4
8854   4
Name: R1, Length: 8478, dtype: int64

```

## Model selection

```

In [40]: df = df.copy()
df["R_score"] = df.mean(axis=1) #row wise mean
df.shape

```

Out[40]: (8478, 9)

```

In [41]: train = df.iloc[:6500]
test = df.iloc[6500:]
train.shape, test.shape

```

Out[41]: ((6500, 9), (1978, 9))

```

In [42]: x_train = train[["R1"]]
y_train = train["R_score"]

```

```
In [43]: model = LinearRegression()
         model.fit(x_train, y_train)
```

```
Out[43]: ▼ LinearRegression ⓘ ?
         LinearRegression()
```

```
In [44]: print("Intercept:", model.intercept_)
         print("Coefficient for R1:", model.coef_[0])

         model.coef_.shape
```

Intercept: 1.018105231431332  
Coefficient for R1: 0.42359934763433976

```
Out[44]: (1,)
```

Hence the form of the best-fit regression line is  $R_{score} = R_1 \times \text{coef} + \text{intercept}$

```
In [45]: y_pred_train = model.predict(x_train)
         RSS_train = np.sum((y_train - y_pred_train) ** 2) #RSS -> residual sum of square
         RSS_train_avg = RSS_train / len(train)
         print("Training RSS (total):", RSS_train)
         print("Training RSS (average):", RSS_train_avg)
```

Training RSS (total): 2902.0393474685015  
Training RSS (average): 0.446467591918231

## Validation

```
In [46]: x_test = test[["R1"]]
         y_test = test[["R_score"]]
```

```
In [47]: y_pred_test = model.predict(x_test)
         RSS_test = np.sum((y_test - y_pred_test) ** 2)
         RSS_test_avg = RSS_test / len(test)
```

```
In [48]: print("Test RSS (total):", RSS_test)
         print("Test RSS (average):", RSS_test_avg)
```

Test RSS (total): 1028.7757850021012  
Test RSS (average): 0.5201090925187569

```
In [49]: print(f"Average RSS (train): {RSS_train_avg:.5f}")
         print(f"Average RSS (test): {RSS_test_avg:.5f}")
```

Average RSS (train): 0.44647  
Average RSS (test): 0.52011

Model works better with training data

## For other fields

### Training

```

In [50]: models = {}
        RSS_train = {}
        RSS_train_avg = {}

        for i in range(2, 9): # R2 to R8
            x_train = train[[f"R{i}"]]
            y_train = train["R_score"]

            models[i] = LinearRegression()
            models[i].fit(x_train, y_train)

            y_pred_train = models[i].predict(x_train)

            RSS_train[i] = np.sum((y_train - y_pred_train) ** 2)
            RSS_train_avg[i] = RSS_train[i] / len(train)

            print(f"R{i}:")
            print(f"  Training RSS (total): {RSS_train[i]}")
            print(f"  Training RSS (average): {RSS_train_avg[i]}\n")

```

R2:  
 Training RSS (total): 2086.9934694791536  
 Training RSS (average): 0.32107591838140825

R3:  
 Training RSS (total): 2856.1295910903655  
 Training RSS (average): 0.43940455247544086

R4:  
 Training RSS (total): 2058.2069482183156  
 Training RSS (average): 0.31664722280281776

R5:  
 Training RSS (total): 2026.8220877210033  
 Training RSS (average): 0.3118187827263082

R6:  
 Training RSS (total): 1762.7887420860695  
 Training RSS (average): 0.2711982680132415

R7:  
 Training RSS (total): 1928.2557723088821  
 Training RSS (average): 0.29665473420136645

R8:  
 Training RSS (total): 1771.9059588635785  
 Training RSS (average): 0.2726009167482428

```

In [51]: dict(sorted(RSS_train_avg.items(), key=lambda item: item[1]))

```

```

Out[51]: {6: np.float64(0.2711982680132415),
          8: np.float64(0.2726009167482428),
          7: np.float64(0.29665473420136645),
          5: np.float64(0.3118187827263082),
          4: np.float64(0.31664722280281776),
          2: np.float64(0.32107591838140825),
          3: np.float64(0.43940455247544086)}

```

```
In [52]: keys = list(RSS_train_avg.keys())
values = list(RSS_train_avg.values())
best_feature = keys[np.argmin(values)]
print(f"Best feature is {best_feature} based on training data")
```

Best feature is 6 based on training data

## Testing

```
In [53]: models = {}
RSS_test = {}
RSS_test_avg = {}

for i in range(2, 9): # R2 to R8
    x_test = test[[f"R{i}"]]
    y_test = test["R_score"]

    models[i] = LinearRegression()
    models[i].fit(x_test, y_test)

    y_pred_test = models[i].predict(x_test)

    RSS_test[i] = np.sum((y_test - y_pred_test) ** 2)
    RSS_test_avg[i] = RSS_test[i] / len(test)

    print(f"R{i}:")
    print(f"  Testing RSS (total): {RSS_test[i]}")
    print(f"  Testing RSS (average): {RSS_test_avg[i]}\n")
```

R2:  
Testing RSS (total): 712.0378480768431  
Testing RSS (average): 0.35997868962428875

R3:  
Testing RSS (total): 875.312355121421  
Testing RSS (average): 0.4425239409107285

R4:  
Testing RSS (total): 799.8393275744875  
Testing RSS (average): 0.4043677085816418

R5:  
Testing RSS (total): 697.1299733893734  
Testing RSS (average): 0.3524418470118167

R6:  
Testing RSS (total): 630.9467824594283  
Testing RSS (average): 0.31898219537888184

R7:  
Testing RSS (total): 680.1640236301748  
Testing RSS (average): 0.3438645215521612

R8:  
Testing RSS (total): 678.3905837840575  
Testing RSS (average): 0.3429679392234871

```
In [54]: dict(sorted(RSS_test_avg.items(), key=lambda item: item[1]))
```

```
Out[54]: {6: np.float64(0.31898219537888184),
          8: np.float64(0.3429679392234871),
          7: np.float64(0.3438645215521612),
          5: np.float64(0.3524418470118167),
          2: np.float64(0.35997868962428875),
          4: np.float64(0.4043677085816418),
          3: np.float64(0.4425239409107285)}
```

```
In [55]: keys = list(RSS_test_avg.keys())
         values = list(RSS_test_avg.values())
         best_feature = keys[np.argmin(values)]
         print(f"Best feature is {best_feature} based on testing data")
```

Best feature is 6 based on testing data

**We can collectively agree that for mean value based regression, 6 is the best performing fitting parameter.**

```
In [ ]:
```